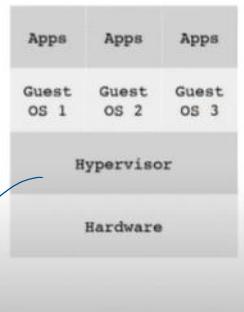


System VM

Hypervisors

- Also called Virtual Machine Monitor (VMM)
- A hypervisor is an operating system for operating systems
 - Provides a virtual execution environment for an entire OS and its applications
 - Controls access to hardware resources
 - When guest OS executes a privileged instruction, Hypervisor intercepts the instruction, checks for correctness and emulates the instruction.



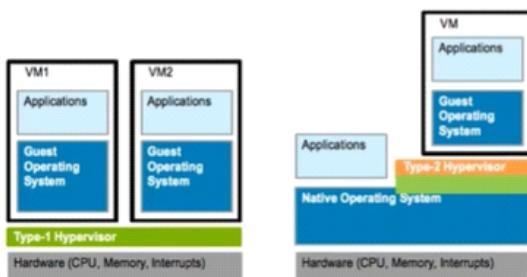
to do this they intercept instructions & emulate it 1st Hypervisor gets booted up guest os does not know hypervisor does this.

does same thing as OS but for VM's.



Type of Hypervisors

- Type 1 Hypervisor (bare metal, native): supports multiple virtual machines and runs directly on the hardware (e.g., VMware ESX, Xen, Denali)
- Type 2 Hypervisor (hosted) VM - runs under a host operating system (e.g., user-mode Linux)



<https://microkerneldude.files.wordpress.com/2012/01/type1-vs-2.png>

VM's will initiate & invoke privileged instructions, but hypervisor will trap them and emulate them.

⇒ Type-2 Hypervisor gives some flexibility

⇒ Type-2 Hypervisor gives some flexibility in terms of implementation.

⇒ we're making Hypervisor part of OS & user space.

⇒ it must've all drivers that is necessary to access hardware.

⇒ why "we write all drivers from scratch"?

- already present in OS.

- why write which Type-1 Hypervisor doing.

⇒ we take advantage of already present drivers.

native host OS — bare metal
guest OS — in VM running

Type-2
Hypervisors

Put Hypervisor in kernel module in OS to take over privileges from host OS & some lib. that are required to run apps.

- Modify guest OS for better performance
- Traditional Hypervisors provide full-virtualization
 - They expose to VMs virtual hardware that is functionally identical to the underlying physical hardware.
 - Advantage : allows unmodified guest OS to execute
 - Disadvantage: Sensitive instructions must be trapped and emulated by Hypervisor.
 - E.g. KVM and VMWare ESX provide full virtualization

- Para-virtualized VM
 - Sees a virtual hardware abstraction that is similar, but not identical to the real hardware.
 - Guest OS is modified to replace sensitive instructions with "hypercalls" to the Hypervisor.
 - Advantage: Results in lower performance overhead
 - Disadvantage: Needs modification to the guest OS.
 - Xen provides both para-virtual as well as full-virtualization

⇒ if Hypervisor wants to get fully transparent
all the instructi. must be trapped & emulated.

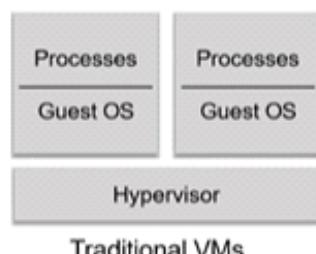
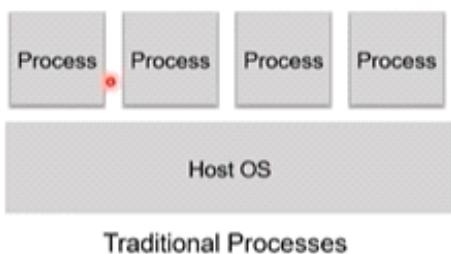
Para-virtualized VMs → guest os modified
not all the instructions
trapped & imulated

- * only sensitive info] one set
- * which "ll be executed through hypervisor
- * & rest instr. "ll be system calls.
- * Hypervisor not fully transparent

- after traditional hypervisors are partially para-virtualized
- device drivers in guest OS may be para-virtualized, whereas CPU & memory may be fully virtualized



Traditional Process vs traditional VMs



Resources shared among processes.

isolation b/w processes that are running on top of these

→ some virtualization / multiplexing done by OS in this level

→ loosely isolated
→ can talk to each other

two guest OS is completely rigid.

→ resources

→ looks very interesting

→ Processes will talk through system calls & e.g. no sys calls in Linux.

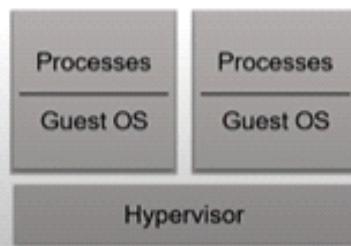
→ all sys. calls available for all research.

→ researchers should be granted as much as possible completely managed by Hypervisor



System VMs

- Each VM has its own
 - Guest OS
 - Guest physical memory ("virtualized" view of memory seen by guest OS)
 - One or more virtual CPUs
 - Virtual I/O devices: virtual disk, virtual network
- Ideally: Co-located VMs don't see/share ANYTHING



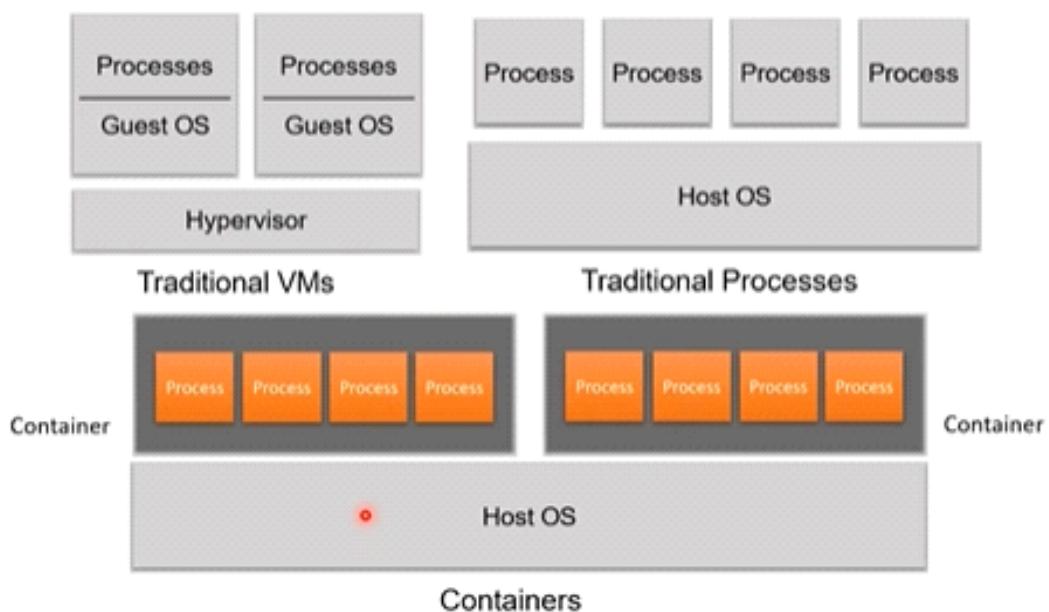


Isolation is important

- Limiting what/who a process/application can see.
- Limiting who can see a process/application
- Processes share too much
 - Great performance but not isolated enough
 - System VMs are too heavy
- Great Isolation but too heavy due to separate guest OS per VM Operating-system-level virtualization
 - Multiple isolated user-spaces
 - Share one kernel
 - Native performance



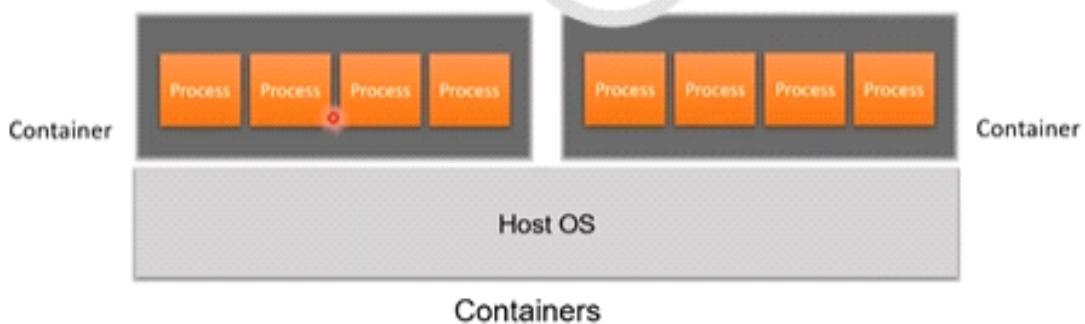
Process VMs and Containers





Containers

- Containers
 - group traditional processes together and
 - restrict what resources they can see/access.
- In Linux, containers consist of
 - Namespaces
 - Control Groups (cgroups)



Processes share namespaces - So if you create one separate namespace & control group of processes .

→ diff. containers will invoke host sys-calls to dominate OS .

→ in terms of sys level APIs a fully virtualized VM - one cannot get full isolation of processes .

But with containers , we can get big reduction approach , while we can've namespaces separated

approach, where we can've namespace separated
for these procedural & isolated reputational sharing
b/w these containers.



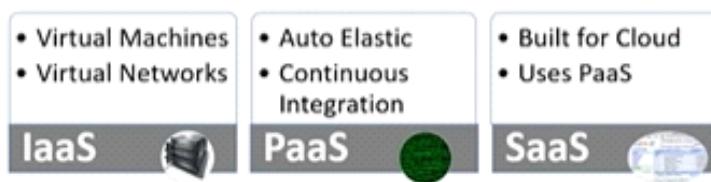
Cloud Computing

- Virtualized distributed processing, storage, and software resources and a service.
- Delivering computing as a on-demand, pay-as-you-go service.
- “A model for enabling convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction” – NIST definition

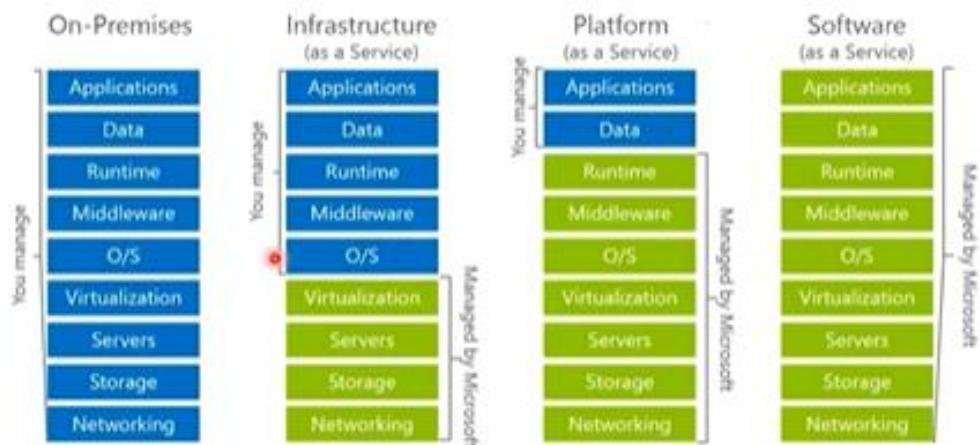


Service models

- IaaS: Infrastructure as a Service
 - Consumer can provision computing resources within provider's infrastructure upon which they can deploy and run arbitrary software, including OS and applications
- PaaS: Platform as a Service
 - Consumer can create custom applications using programming tools supported by the provider and deploy them onto the provider's cloud infrastructure
- SaaS: Software as a Service
 - Consumer uses provider's applications running on provider's cloud infrastructure



Service models



IAAS cloud acrhitecture

