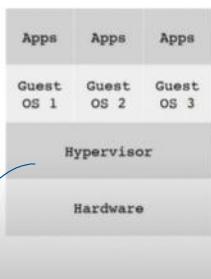


## System VM

### Hypervisors

- Also called Virtual Machine Monitor (VMM)
- A hypervisor is an operating system for operating systems
  - Provides a virtual execution environment for an entire OS and its applications
  - Controls access to hardware resources
  - When guest OS executes a privileged instruction, Hypervisor intercepts the instruction, checks for correctness and emulates the instruction.

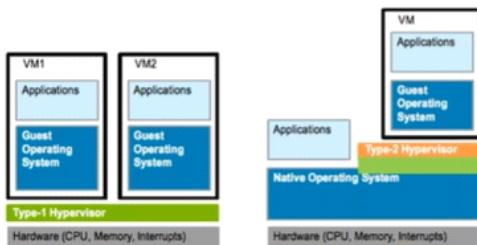


to do this they intercept instructions & emulates if guest os does not know hypervisor does this. 1<sup>st</sup> Hypervisor gets booted up does same thing as OS but for VM's



## Type of Hypervisors

- Type 1 Hypervisor (bare metal, native): supports multiple virtual machines and runs directly on the hardware (e.g., VMware ESX, Xen, Denali)
- Type 2 Hypervisor (hosted) VM - runs under a host operating system (e.g., user-mode Linux)



<https://microkerneldude.files.wordpress.com/2012/01/type1-vs-2.png>

VM's will initiate & invoke privileged instructions, but hypervisor will trap them and emulate them.

⇒ Type-2 Hypervisor gives some flexibility in terms of implementation.

⇒ we're making hypervisor part of OS & user space.  
↳ ... via all drivers that is necessary to access

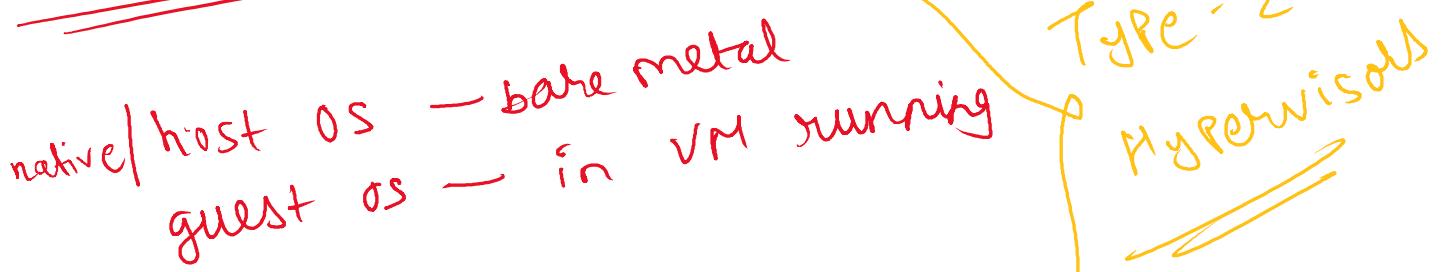


⇒ We're making Hypervisor part of OS & now we  
⇒ It must've all drivers that is necessary to access  
hardware.

⇒ Why'll write all drivers from scratch?

- already present in OS.
- Why write which Type-1 Hypervisor doing.

# We take advantage of already present drivers.



Put Hypervisor in kernel  
module in OS to take over  
privileged from host OS &  
some lib. that are required  
to run apps.



## Para-virtualized VMs



- Modify guest OS for better performance
- Traditional Hypervisors provide full-virtualization
  - They expose to VMs virtual hardware that is functionally identical to the underlying physical hardware.
  - Advantage: allows unmodified guest OS to execute
  - Disadvantage: Sensitive instructions must be trapped and emulated by Hypervisor.
  - E.g. KVM and VMWare ESX provide full virtualization
- Para-virtualized VM
  - Sees a virtual hardware abstraction that is similar, but not identical to the real hardware.
  - Guest OS is modified to replace sensitive instructions with "hypercalls" to the Hypervisor.
  - Advantage: Results in lower performance overhead



- Para-virtualized VM
  - Sees a virtual hardware abstraction that is similar, but not identical to the real hardware.
  - Guest OS is modified to replace sensitive instructions with "hypervcalls" to the Hypervisor.
  - Advantage: Results in lower performance overhead
  - Disadvantage: Needs modification to the guest OS.
  - Xen provides both para-virtual as well as full-virtualization

→ if Hypervisor wants to get fully transparent all the instructi. must be trapped & emulated.



guest os modified  
not all the instruction  
trapped & immediate

# only sensitive  
info. ] one set

# which -> ve executed through hyp

# & rest instr. -> be system co

# Hypervisor not fully transpar

→ often + additional hy  
i...ll...

permissive  
cells.  
agent

permissive  
inhalized

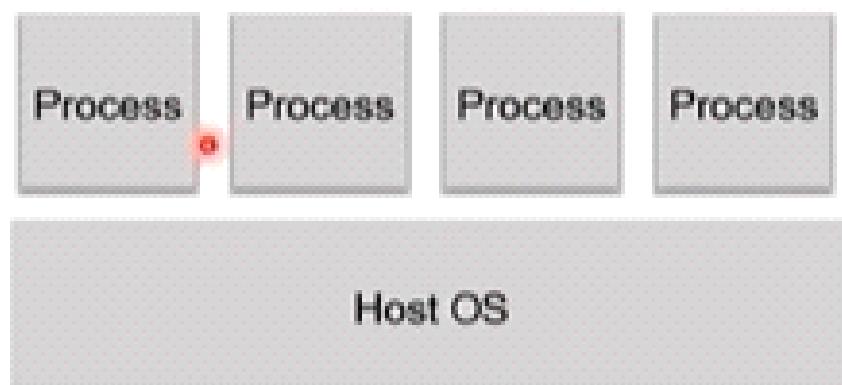
are partially Pa

device drivers in g  
Para-virtualized  
CPU & memory ma

partially virtualized  
host OS may be  
, whereas  
they be fully virtualized



# Traditional Process vs traditional

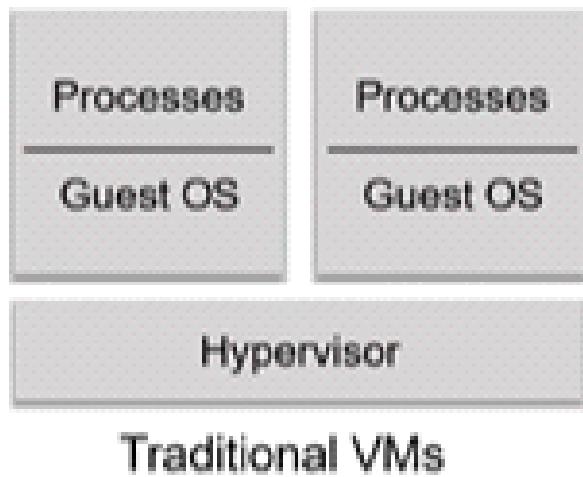


Traditional Processes

Resources shared  
among processes.

→ some virtualization / multi

# Virtual VMs



isolation b/w processes  
that are running  
on top of these  
two guest os is  
completely  
plexing

- some virtualization / memory management done by OS in this layer
- loosely isolated
- Processes will talk through system calls & Eg. sys calls in Linux.
- all sys. calls available resources.

Completely  
rigid.

→ resources  
shared b/w  
guest os's  
completely  
managed by  
Hypervisor

through  
400  
for all



## System VMs

- Each VM has its own
  - Guest OS
  - Guest physical memory ("virtualized" view of memory)
  - One or more virtual CPUs
  - Virtual I/O devices: virtual disk, virtual network interface card
- Ideally: Co-located VMs don't see/share ANYTHING



---

memory seen by guest OS'

2



# Isolation is important

- Limiting what/who a process/application can see.
- Limiting who can see a process/application
  
- Processes share too much
  - Great performance but not isolated enough
  - System VMs are too heavy
- Great Isolation but too heavy due to separate guest system-level virtualization
  - Multiple isolated user-spaces
  - Share one kernel
  - Native performance

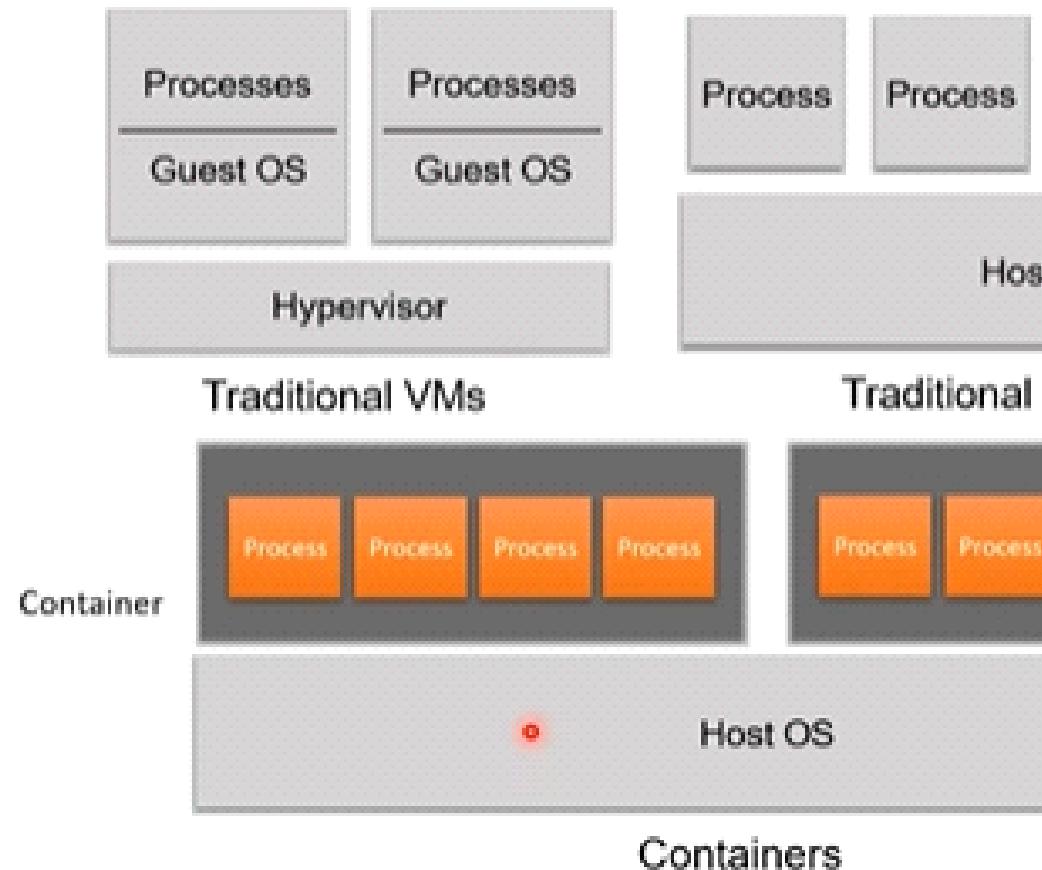
---

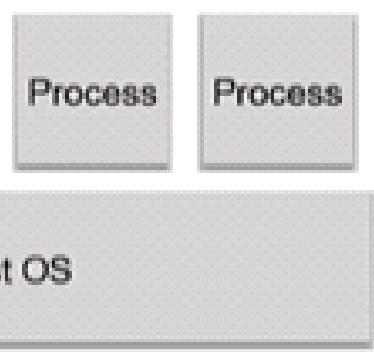
et OS per VM Operating-



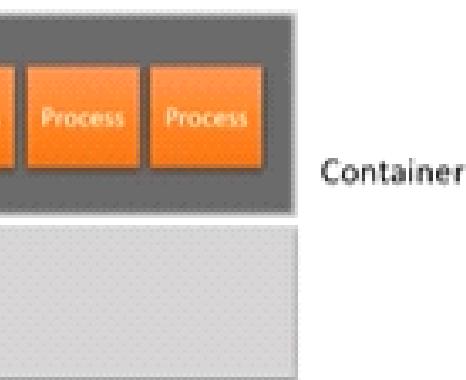


# Process VMs and Containers





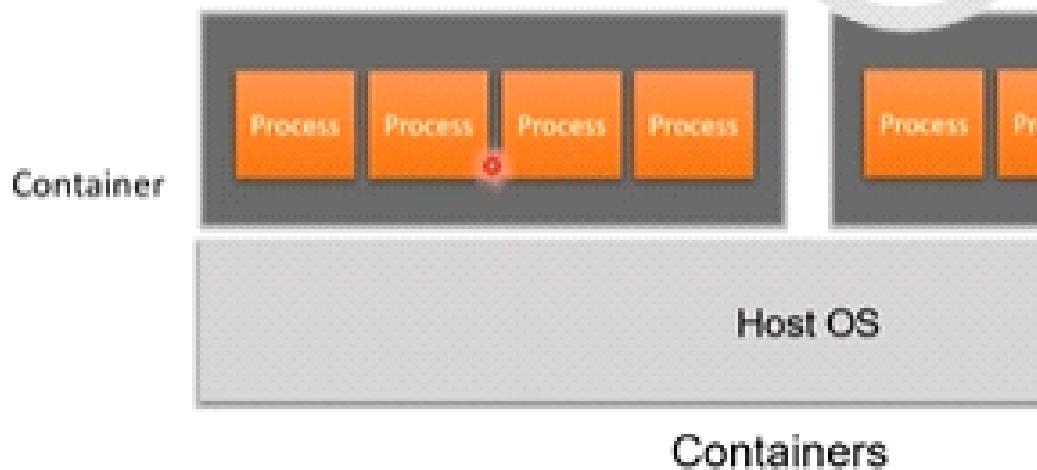
## Processes



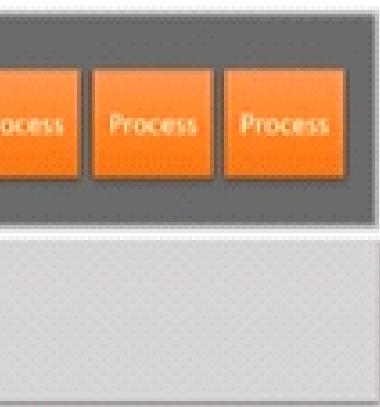


# Containers

- Containers
  - group traditional processes together and
  - restrict what resources they can see/access
- In Linux, containers consist of
  - Namespaces
  - Control Groups (cgroups)



Processes share names  
create one separate  
group of processes .



spaces, So if you  
namespace & control

⇒ diff. containers will invoke  
to resource OS.

⇒ in terms of sys level v.  
one cannot get full  
But with container , we can  
approach , where we can  
run these processes . &  
blw these containers -

invoke host sys. calls

$M_s$  or fully virtualized  $VM_s$

isolation of processes.

We can get big resource

an've namespace separated

isolated resource sharing



# Cloud Computing

- Virtualized distributed processing, storage service.
- Delivering computing as a on-demand, pay
- "A model for enabling convenient, on-demand pool of configurable computing resources (e.g. networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider definition"

, and software resources and a

y-as-you-go service.

band network access to a shared  
(e.g., networks, servers, storage,  
idly provisioned and released with  
ovider interaction” – NIST



# Service models

- IaaS: Infrastructure as a Service
  - Consumer can provision computing resources upon which they can deploy and run arbitrary applications
- PaaS: Platform as a Service
  - Consumer can create custom applications by the provider and deploy them onto the provider's infrastructure
- SaaS: Software as a Service
  - Consumer uses provider's applications running on infrastructure

- Virtual Machines
- Virtual Networks

IaaS

- Auto Elastic
- Continuous Integration

PaaS

ces within provider's infrastructure  
ary software, including OS and

using programming tools supported  
provider's cloud infrastrcture

nning on provider's cloud

- Built for Cloud
- Uses PaaS

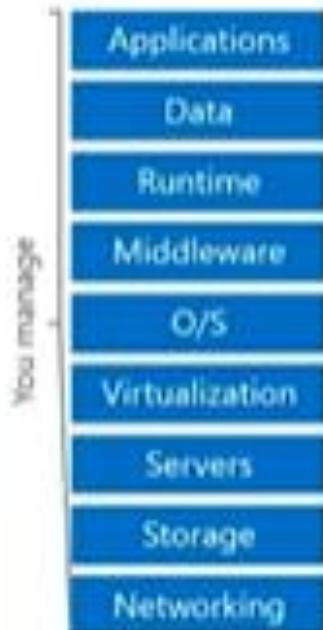
SaaS



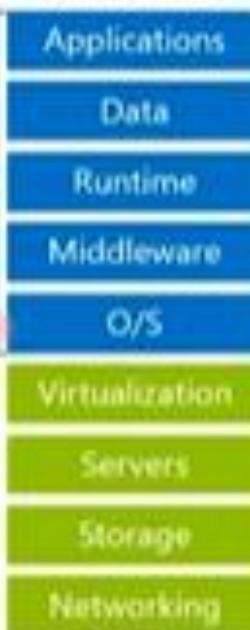


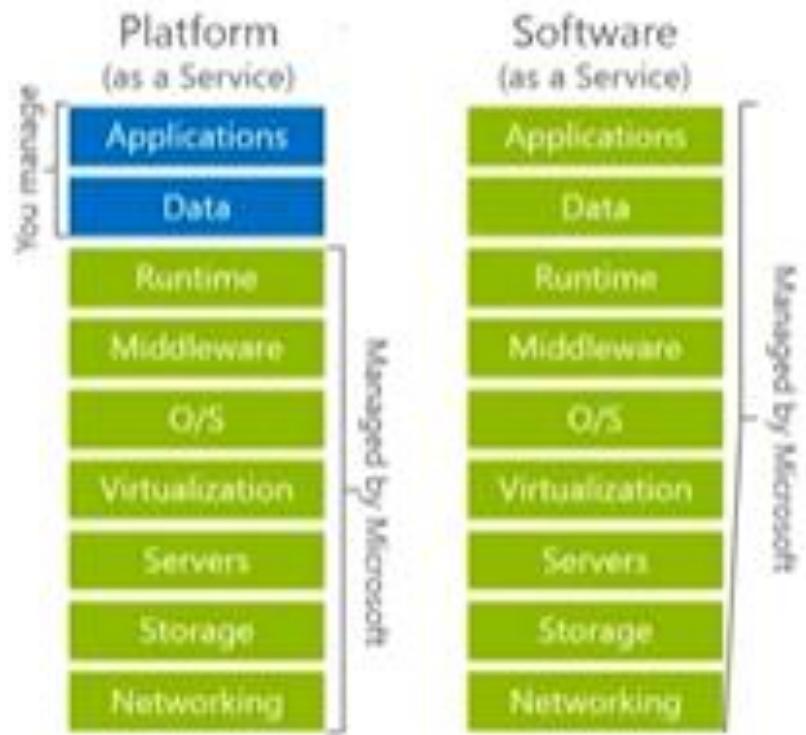
# Service models

## On-Premises



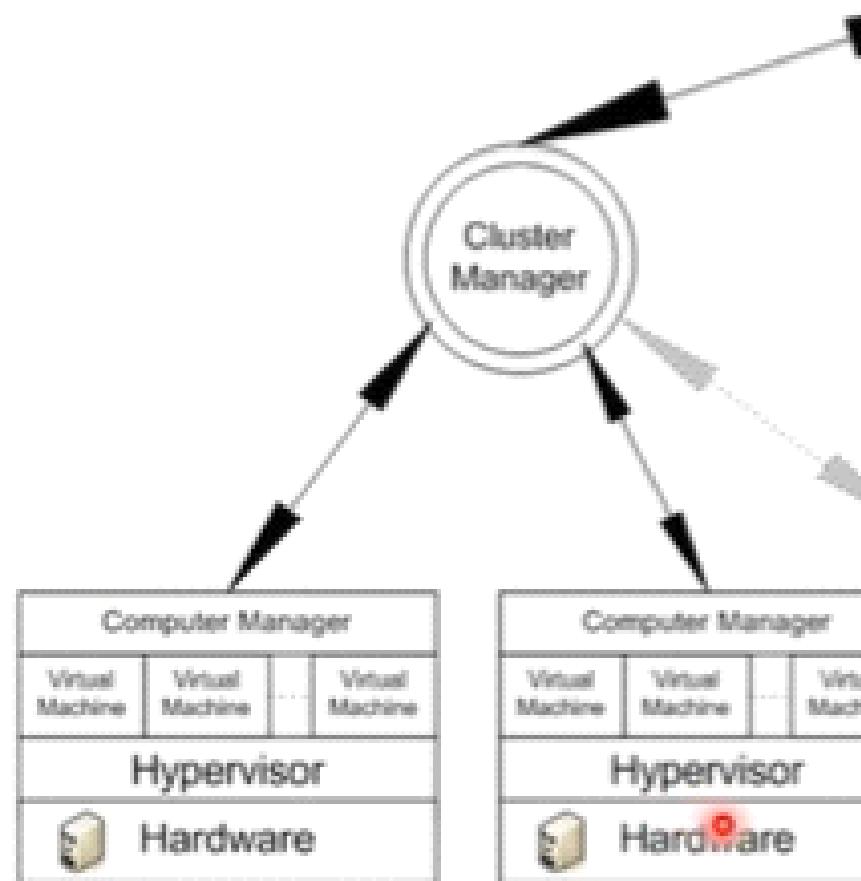
## Infrastructure (as a Service)





# IAAS cloud acrhitectur

IaaS Architecture



re

